

DCP: Addressing Input Dynamism In Long-Context Training via Dynamic Context Parallelism


Chenyu Jiang, Zhenkun Cai, Ye Tian, Zhen Jia, Yida Wang, Chuan Wu

The University of Hong Kong, Amazon Web Services




Background

Context Parallelism




Background

Context Parallelism



Background



Context Parallelism




Setup: 8B model, Megatron-LM, 400Gbps interconnect between nodes, 8/16-way context parallelism cross nodes

Motivation

Input Dynamism *in sequence length*




FLANv2



The Pile


Motivation

Input Dynamism *in sequence length*




Motivation

Input Dynamism *in sequence length*





(full scale communication required for all sequences)




Motivation

Input Dynamism *in sequence length*




Motivation

Input Dynamism *in attention masks*




StreamingLLM

Xiao, et al., 2024



Causal Blockwise Mask

Bertsch, et al., 2025





Shared Question Mask

Wang, et al., 2025

Diverse attention mask patterns on input sequences.

Motivation



Input Dynamism *in attention masks*



Special Data Placement for Causal Mask

Motivation

Input Dynamism *in attention masks*




Special Data Placement for Causal Mask


Motivation

Input Dynamism *in attention masks*


Other masks...



StreamingLLM




Causal Blockwise Mask




Shared Question Mask

Motivation

Input Dynamism *in attention masks*




Computation Imbalance




Motivation

Input Dynamism *in attention masks*



Transfer of KV



Redundant Communication

Motivation

Input Dynamism *in distributed attention*

Observation:

Different batches require **different parallelization strategy** for optimal performance.

Question:

How to automatically optimize parallelization strategy for each input batch?

How to build a system that flexibly adapts to such dynamic parallelism?

Design

Path to automatic parallelization strategy optimization

1. Optimize the placement of data and computation (parallelization)
2. Determine the schedule of communication and computation

Design

Path to automatic parallelization strategy optimization


1. Optimize the placement of data and computation (parallelization)

2. Determine the schedule of communication and computation

Design

Optimize parallelization with Hypergraph Partition


Attention described with two type of blocks:



Design

Optimize parallelization with Hypergraph Partition


Attention described with two type of blocks:



Design

Optimize parallelization with Hypergraph Partition


Attention described with two type of blocks:



Design

Optimize parallelization with Hypergraph Partition

Attention described with two type of blocks:



Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication


Hypergraph: each edge connects ≥ 2 vertices

Vertices: data and computation blocks

Hyperedges: dependency between vertices
(one for each data block)

Communication: for each hyper-edge, required communication is:

$$(\# \text{cut} - 1) \times \text{Size(data block)}$$



Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication

Hypergraph: each edge connects ≥ 2 vertices


Vertices: data and computation blocks

Hyperedges: dependency between vertices

(one for each data block)

Communication: for each hyper-edge, required communication is:

$$(\# \text{ cut} - 1) \times \text{Size(data block)}$$



Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication


Hypergraph: each edge connects ≥ 2 vertices

Vertices: data and computation blocks

Hyperedges: dependency between vertices
(one for each data block)

Communication: for each hyper-edge, required communication is:

$$(\# \text{cut} - 1) \times \text{Size(data block)}$$



Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication

Hypergraph: each edge connects ≥ 2 vertices


Vertices: data and computation blocks

Hyperedges: dependency between vertices

(one for each data block)

Communication: for each hyper-edge, required communication is:

$$(\# \text{ cut} - 1) \times \text{Size(data block)}$$



Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication


Hypergraph: each edge connects ≥ 2 vertices

Vertices: data and computation blocks

Hyperedges: dependency between vertices
(one for each data block)

Communication: for each hyper-edge, required communication is:

$$(\# \text{cut} - 1) \times \text{Size(data block)}$$




Design

Optimize parallelization with Hypergraph Partition

Goal: find a **balanced** partition of data and computation, while **minimizing** communication

Solving the balanced hyper-graph partitioning problem yields the optimal data and computation placement.



Design


Path to automatic parallelism optimization

1. Optimize the placement of data and computation (parallelization)
2. Determine the schedule of communication and computation

Design

Block scheduling for overlapping computation and communication

Goal: maximize communication-computation **overlap** while avoiding **congestion**



Design



Block scheduling for overlapping computation and communication

Partition comm. and comp. on each device into **divisions**.

Within each division, desire **balanced computation and communication**.


Communication required by the next division can **overlap** with computation in the current division.

Problem is **NP-hard**, using a greedy heuristic.



Design

System Overview




```
# when defining models
class TransformerLayer(...):
    def forward(..., dcp_executor):
        ...
        # replace attention implementation with DCPOpt
        core_attn_out = DCPOpt.apply(dcp_executor, q, kv)
        ...

# in training script
dcp_dataloader = DCPDataloader(dataset, mask_fn)
# dcp_group is a communicator that connects all devices
# (e.g., torch.distributed.ProcessGroup)
dcp_executor = DCPExecutor(group=dcp_group)
# training iterations
for (local_data, execution_plan) in dcp_dataloader:
    # set execution plan and create buffers
    dcp_executor.prepare(execution_plan)
    # execute model
    loss = model(local_data, dcp_executor)
    ...
```

Design

System Overview




```
# when defining models
class TransformerLayer(...):
    def forward(..., dcp_executor):
        ...
        # replace attention implementation with DCPAttn
        core_attn_out = DCPAttn.apply(dcp_executor, q, kv)
        ...

# in training script
dcp_dataloader = DCPDataloader(dataset, mask_fn)
# dcp_group is a communicator that connects all devices
# (e.g., torch.distributed.ProcessGroup)
dcp_executor = DCPExecutor(group=dcp_group)
# training iterations
for (local_data, execution_plan) in dcp_dataloader:
    # set execution plan and create buffers
    dcp_executor.prepare(execution_plan)
    # execute model
    loss = model(local_data, dcp_executor)
    ...
```

Design

System Overview




```
# when defining models
class TransformerLayer(...):
    def forward(..., dcp_executor):
        ...
        # replace attention implementation with DCPOpt
        core_attn_out = DCPOpt.apply(dcp_executor, q, kv)
        ...

# in training script
dcp_dataloader = DCPOptDataloader(dataset, mask_fn)
# dcp_group is a communicator that connects all devices
# (e.g., torch.distributed.ProcessGroup)
dcp_executor = DCPEExecutor(group=dcp_group)
# training iterations
for (local_data, execution_plan) in dcp_dataloader:
    # set execution plan and create buffers
    dcp_executor.prepare(execution_plan)
    # execute model
    loss = model(local_data, dcp_executor)
    ...
```

Design

System Overview



```
# when defining models
class TransformerLayer(...):
    def forward(..., dcp_executor):
        ...
        # replace attention implementation with DCPOpt
        core_attn_out = DCPOpt.apply(dcp_executor, q, kv)
        ...


# in training script
dcp_dataloader = DCPOptDataloader(dataset, mask_fn)
# dcp_group is a communicator that connects all devices
# (e.g., torch.distributed.ProcessGroup)
dcp_executor = DCPExecutor(group=dcp_group)

# training iterations
for (local_data, execution_plan) in dcp_dataloader:
    # set execution plan and create buffers
    dcp_executor.prepare(execution_plan)
    # execute model
    loss = model(local_data, dcp_executor)
    ...
```





Key Results

Configuration: 8 nodes, each with 8 A100-80GB GPUs, 400 Gbps interconnect between nodes





Model: follow llama3-8B setup



Attention microbenchmark: speed up 1.19x~2.45x under causal masks, 2.15x~3.77x under sparse masks




End-to-end training performance on the LongAlign dataset.




End-to-end training performance on the LongDataCollections dataset.

End-to-end benchmarks: speed up 0.94x~1.16x under causal masks, 1.00x~1.46x under sparse masks


Key Results



LongAlign.



LongDataCollections.



Caption

Planning time: < 10s per iteration under reasonable block size, full overlap with model execution when parallelized onto more than 10 CPU cores.

Timeline decomposition: communication time greatly reduced. Potential performance improvement with better communication scheduling.

Takeaway

Dynamism in model input — sequence length and attention mask — can be exploited to accelerate context parallelism training.

Context parallelization strategies can be described and optimized in two layers:

1. data and computation placement (modelled and optimized with hyper-graph representation)
2. Computation/communication scheduling.

Thank you



Contact: jchenyu@connect.hku.hk