

# DynaPipe: Optimizing Multi-task Training through Dynamic Pipelines

Chenyu Jiang<sup>1\*</sup>, Zhen Jia<sup>2</sup>, Shuai Zheng<sup>2</sup>, Yida Wang<sup>2</sup>, Chuan Wu<sup>1</sup>

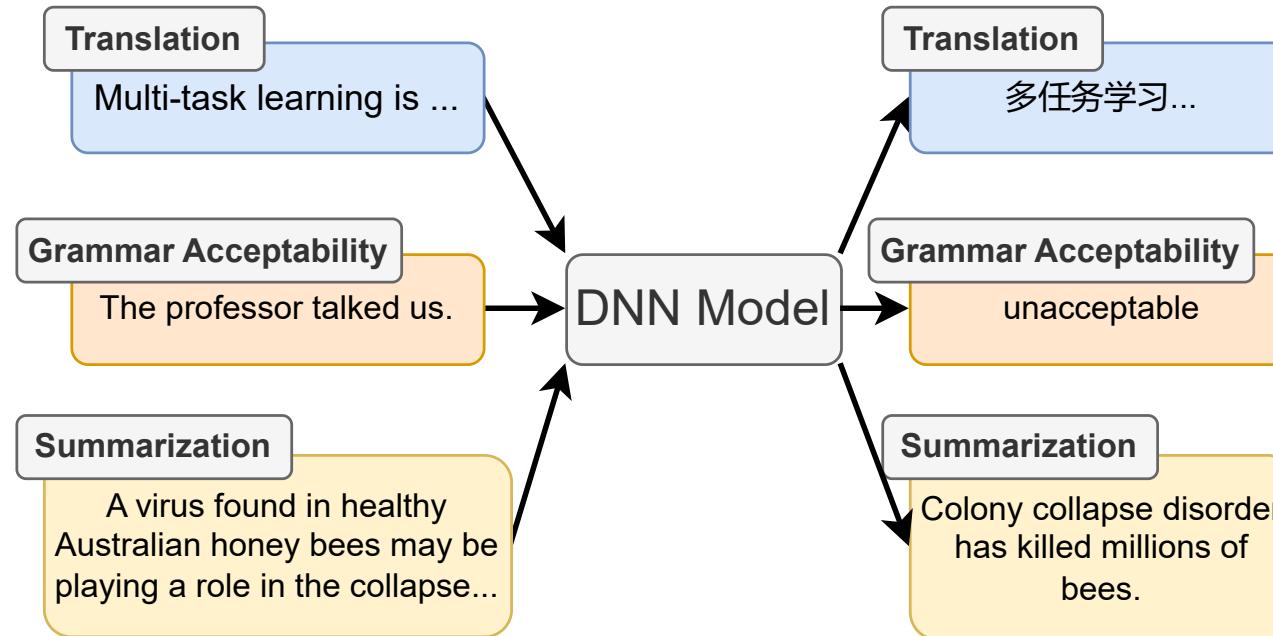
<sup>1</sup>The University of Hong Kong, <sup>2</sup>Amazon Web Services

Accepted to EuroSys 2024

\*Work done while interning at AWS.

# Background

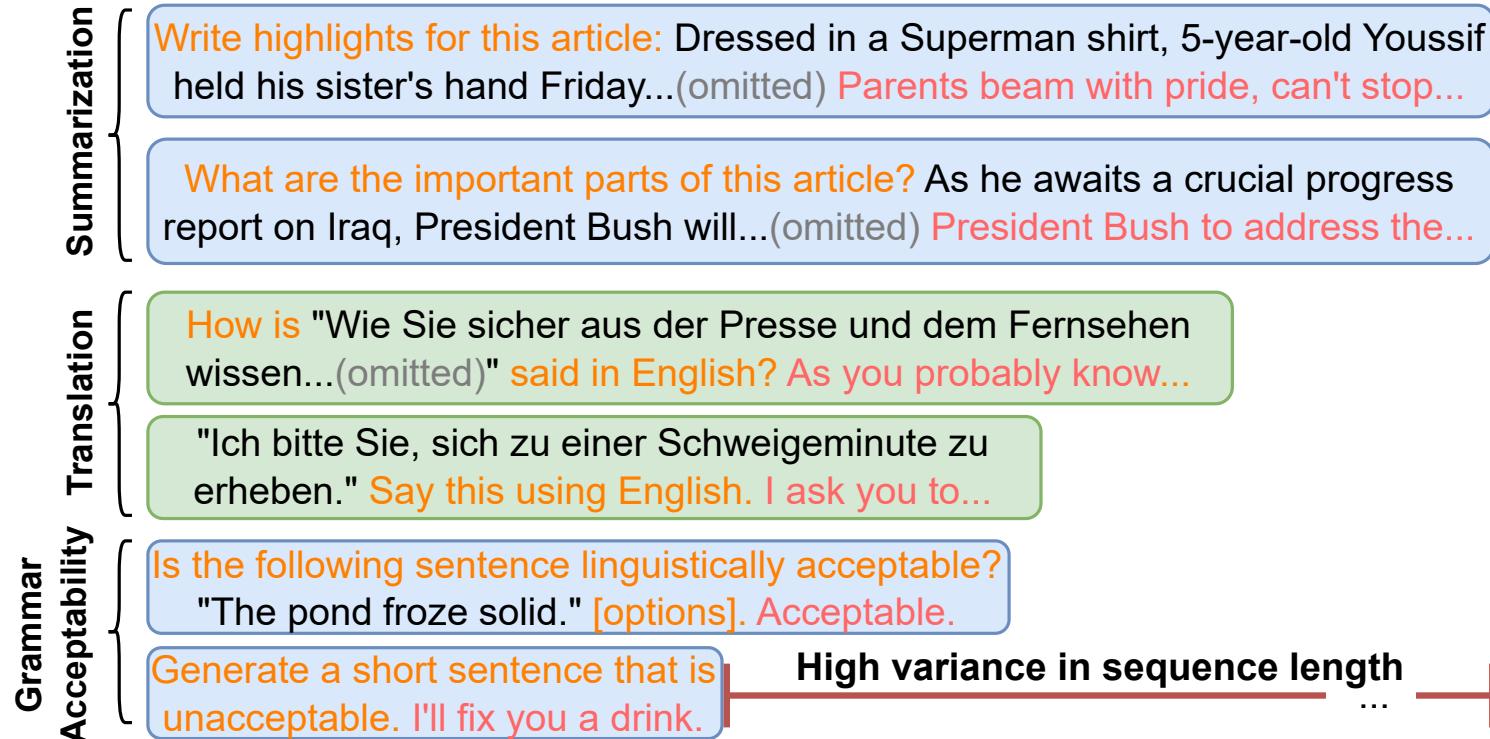
## Multi-task Training



One DNN model, multiple tasks.

# Background

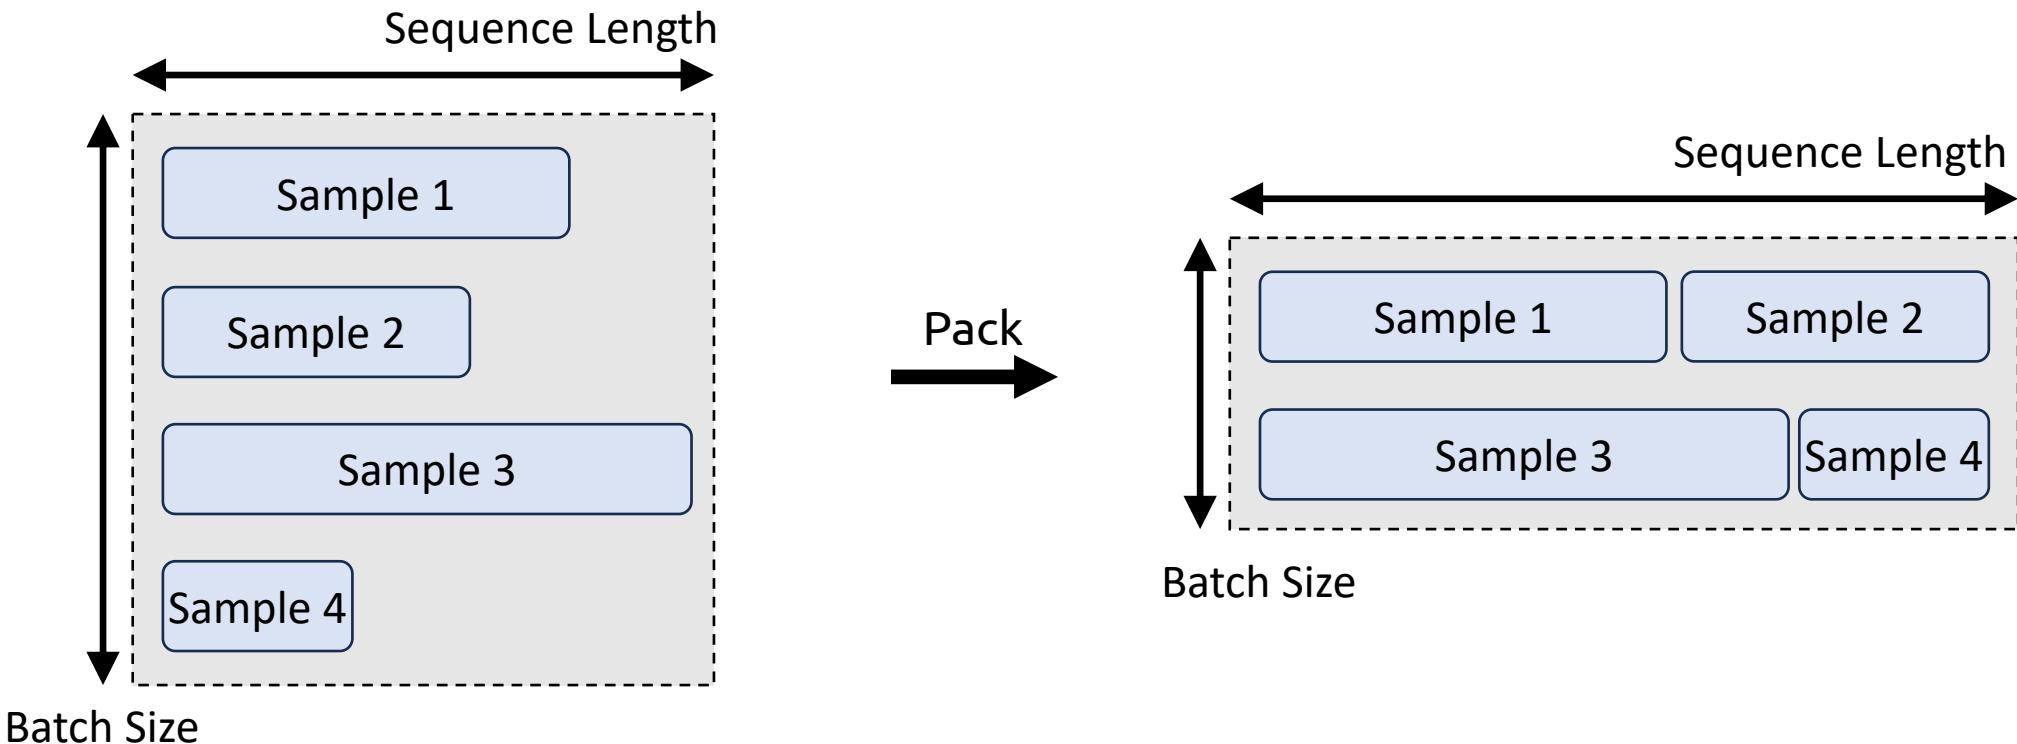
## High sequence length variation in multi-task training



Large amount of padding needed.

# Background

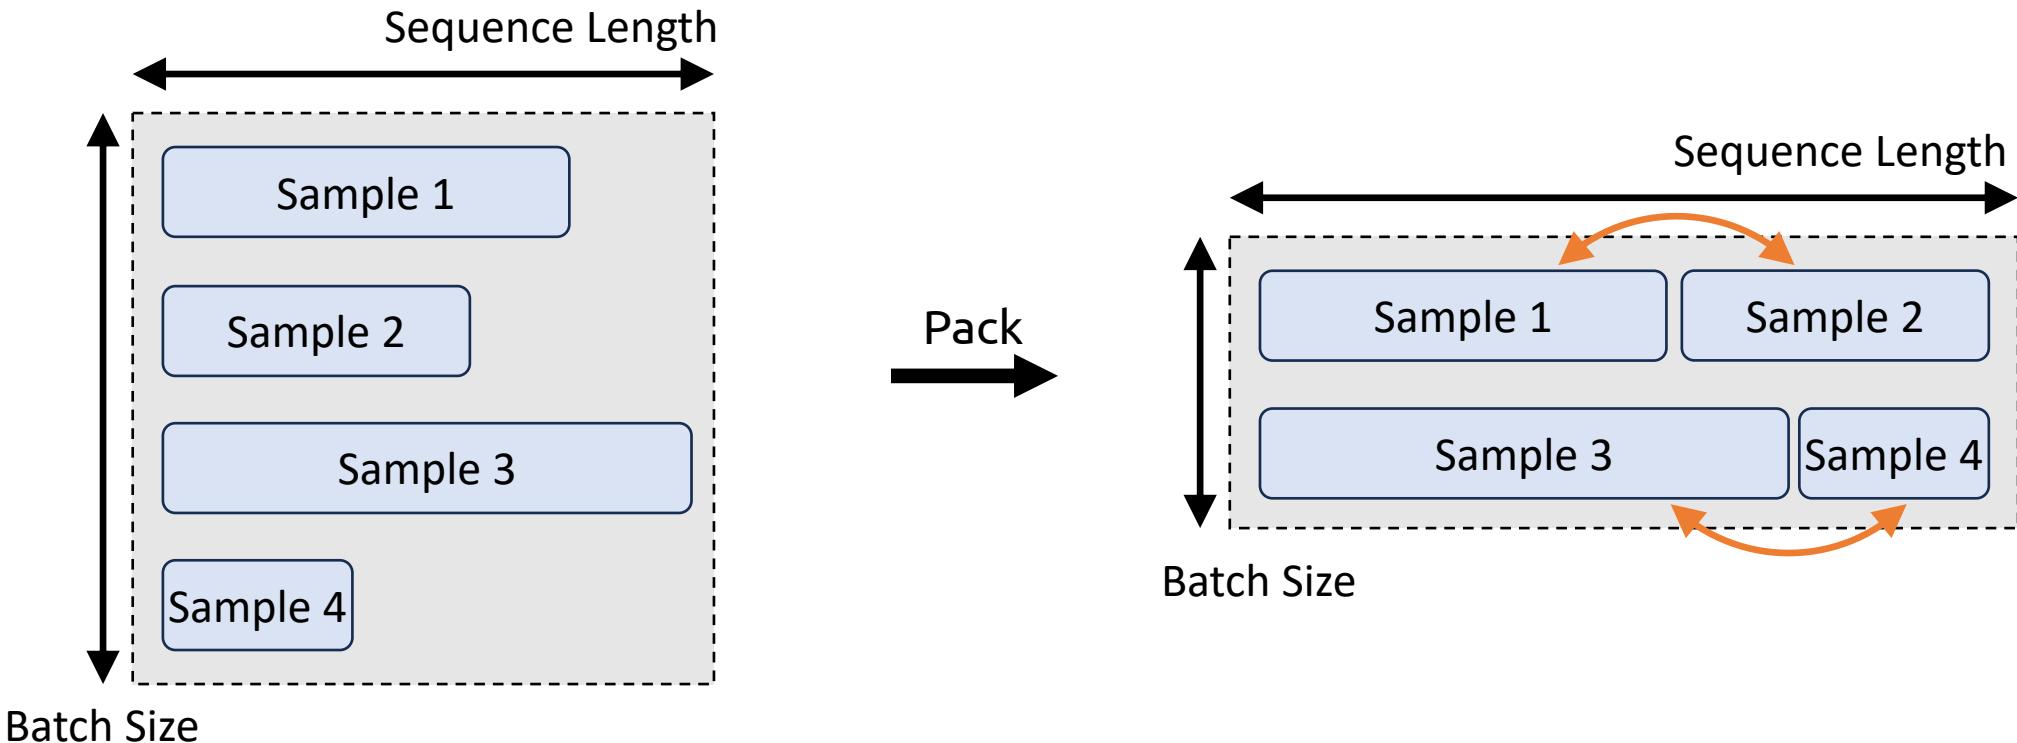
Current solution: packing



Concatenating short sequences to long ones up to fixed maximum length.

# Background

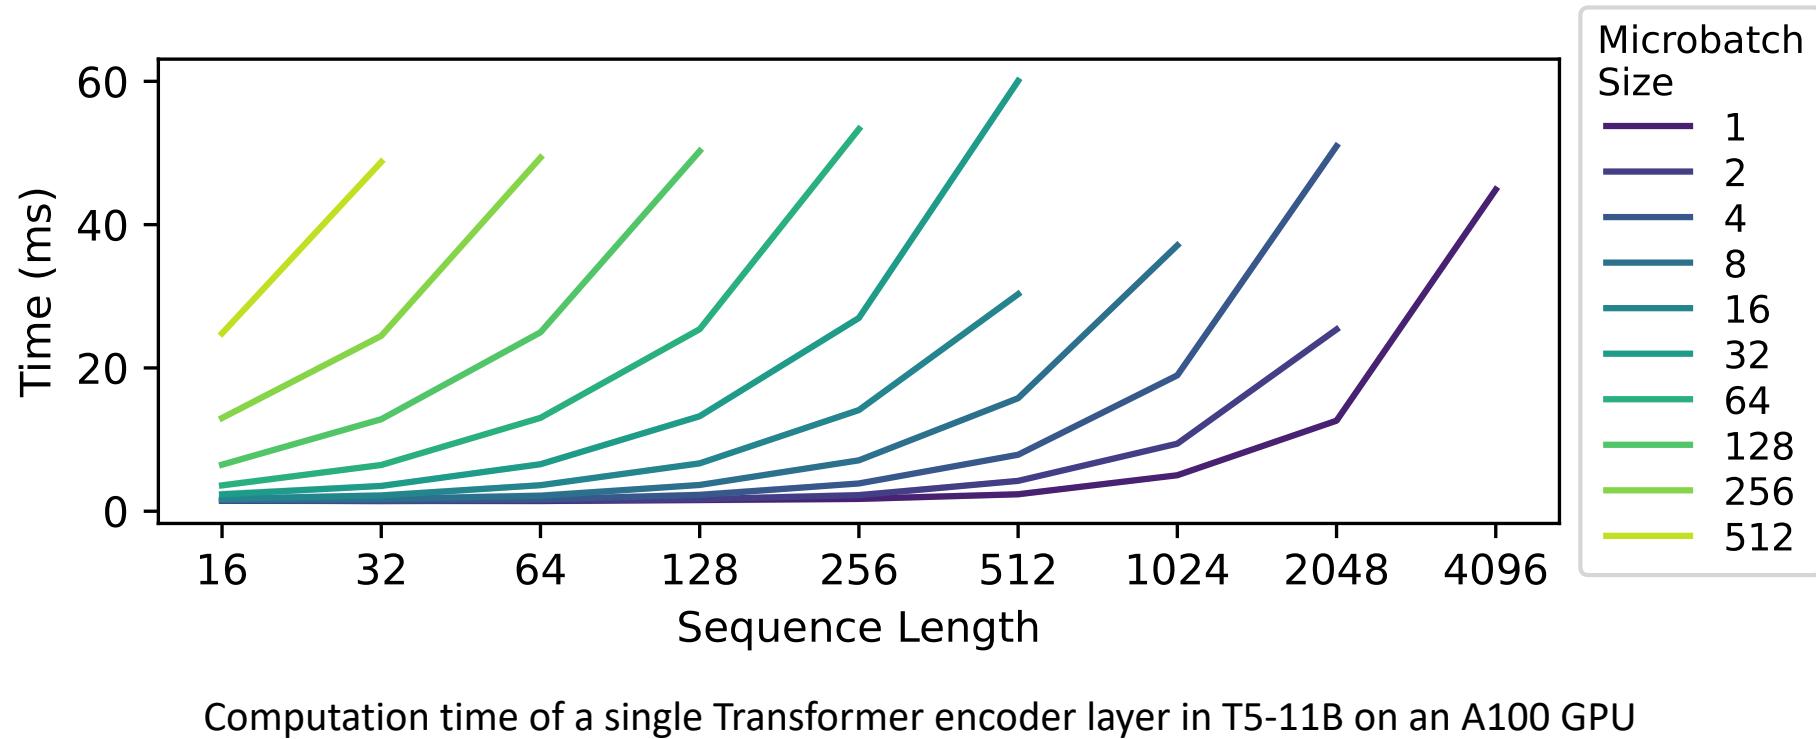
Current solution: packing



Drawback: **unnecessary attention** between unrelated samples.

# Background

Current solution: packing

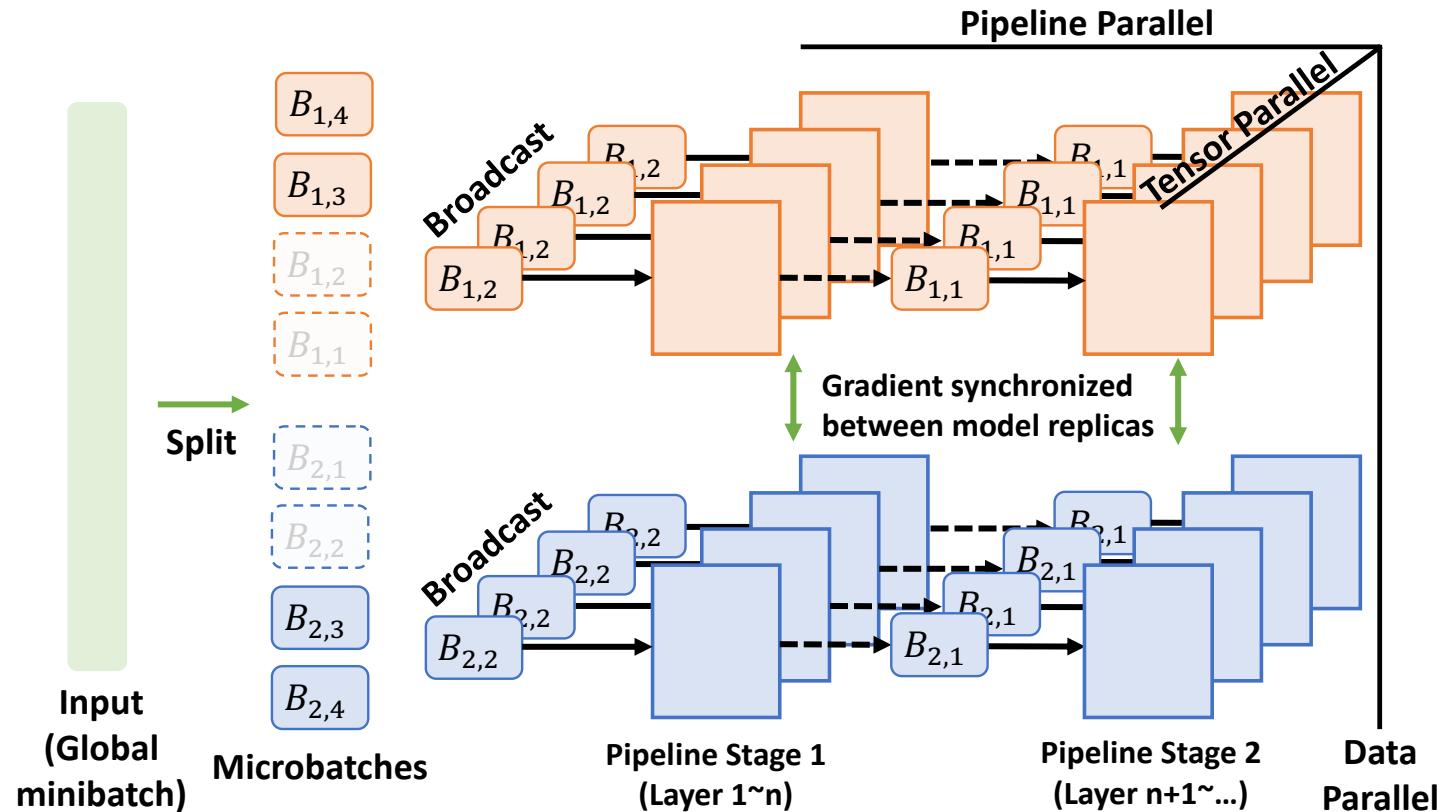


Drawback: **super-linear** execution time growth with sequence length

# Motivation

Proposed solution: dynamic micro-batching

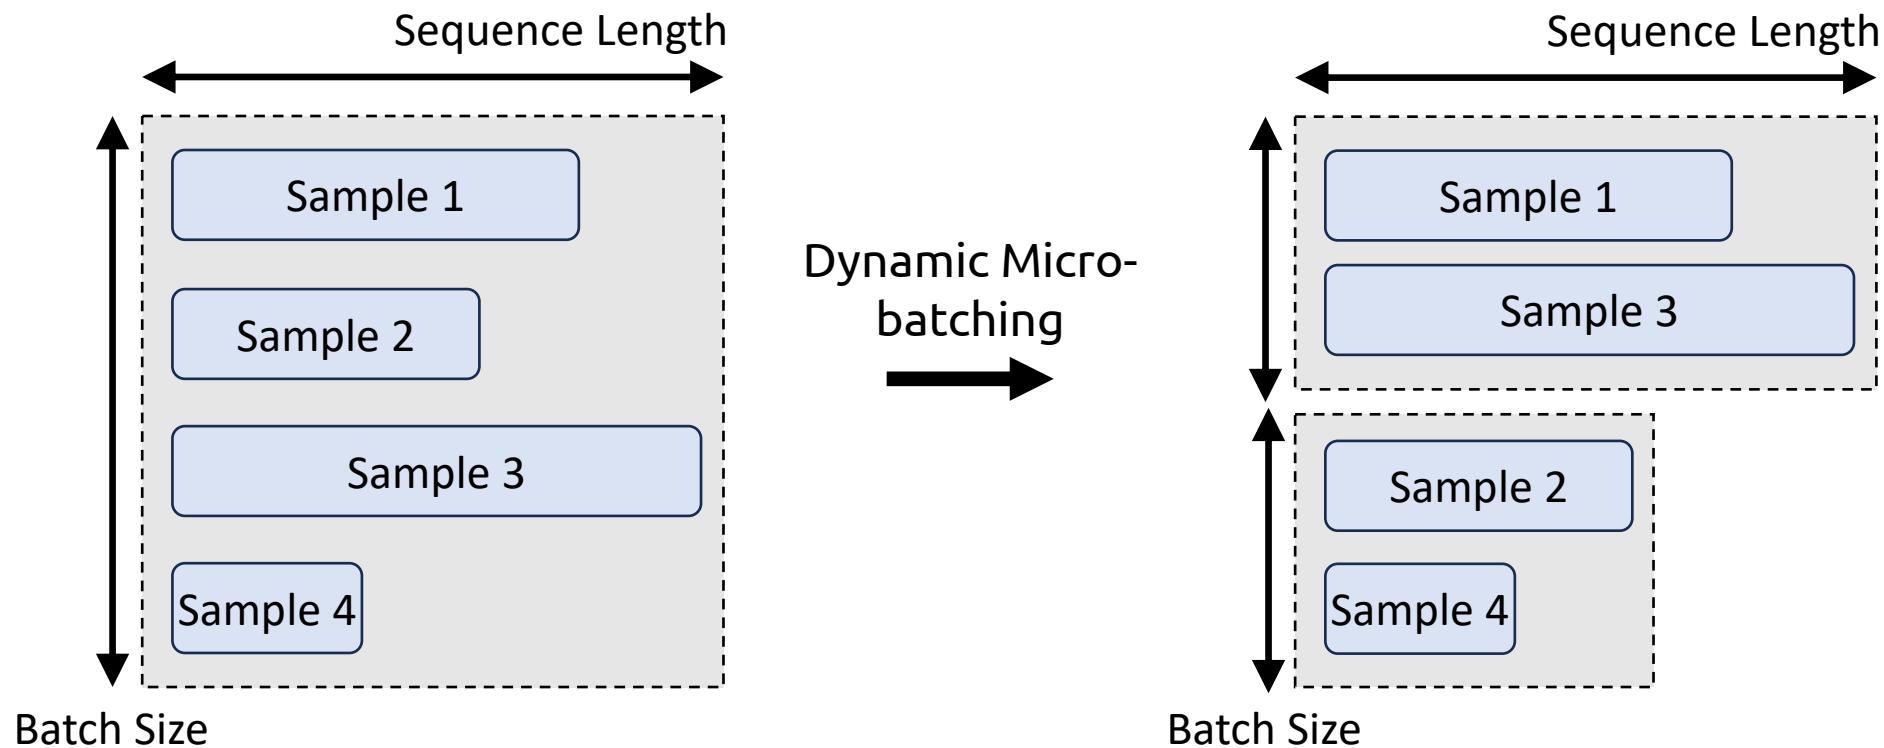
Optimally split each input global batch into micro-batches with similar sequence lengths.



# Motivation

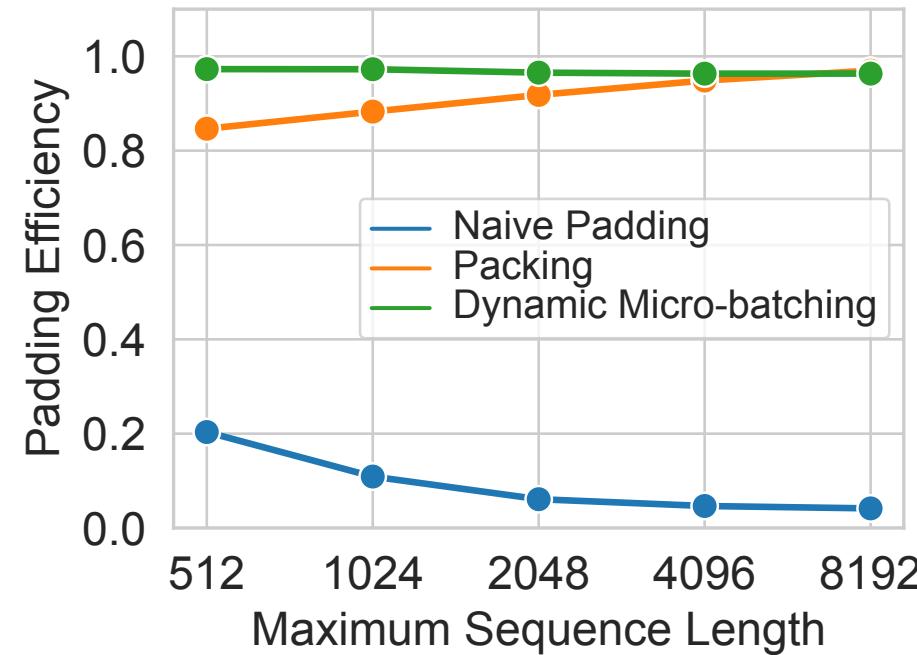
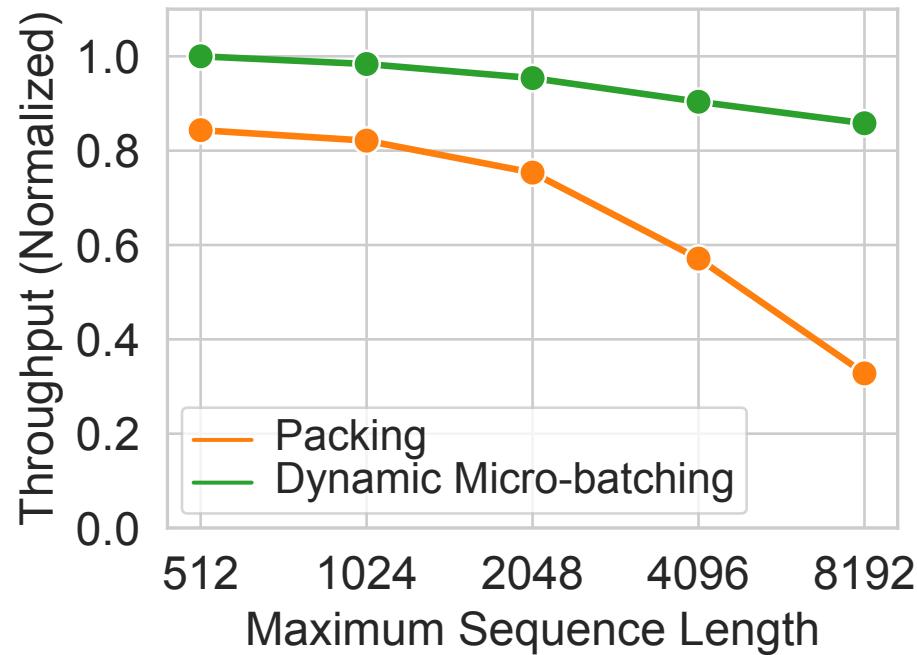
Proposed solution: dynamic micro-batching

Optimally split each input global batch into micro-batches with similar sequence lengths.



# Motivation

Proposed solution: dynamic micro-batching

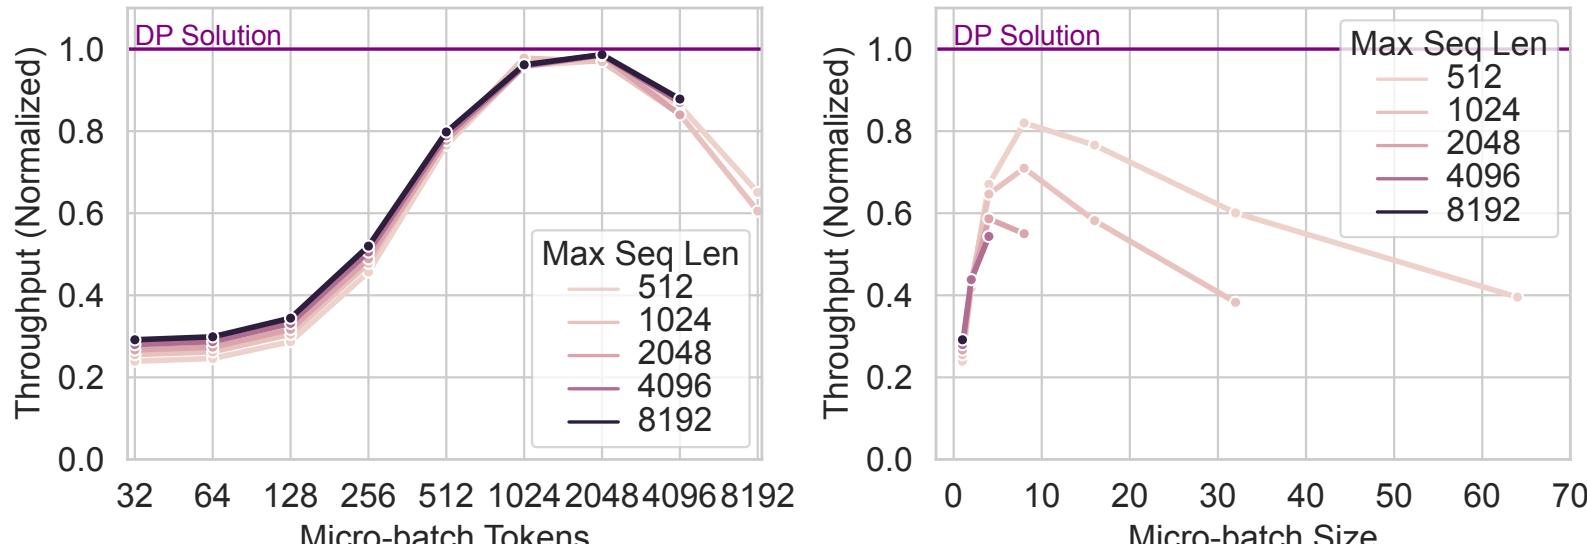


Comparison of throughput and padding efficiency of different methods when training a GPT model with FLANv2 dataset

# Motivation

## Challenges of dynamic micro-batching

**No principled way to split training mini-batches into micro-batches of different sequence lengths.**



Training performance of GPT under different micro-batching methods

Micro-batching is critical for performance, yet difficult to find the best batching configuration.

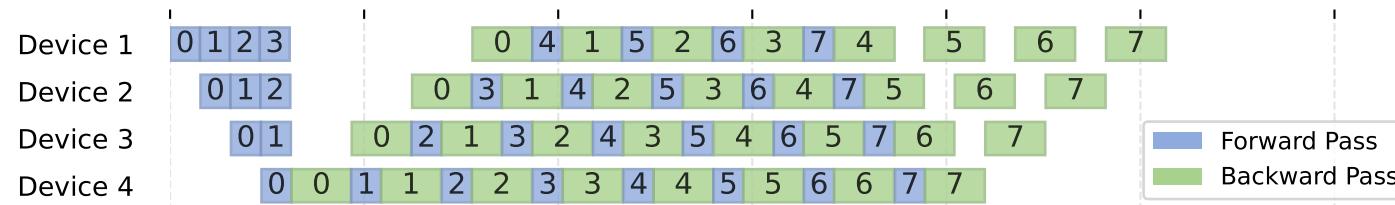
Our solution: a dynamic programming algorithm to optimally construct micro-batches

# Motivation

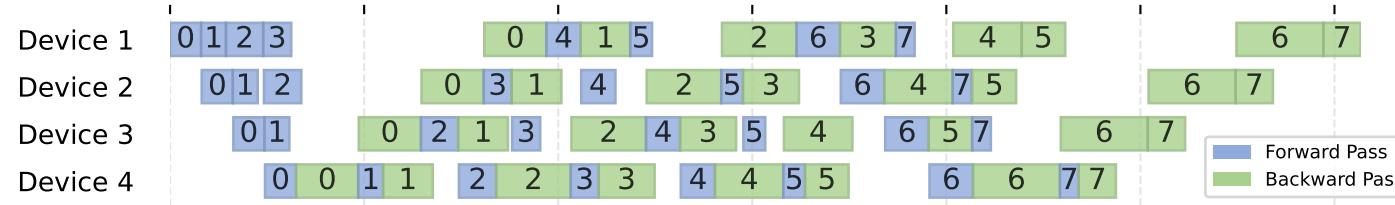
# Challenges of dynamic micro-batching

No efficient pipeline schedules for micro-batches of diverse execution times.

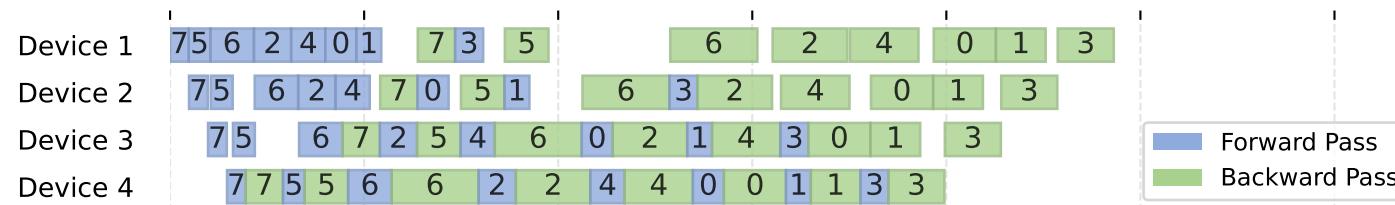
## 1F1B (uniform micro-batches)



## 1F1B (dynamic micro-batches)



# DynaPipe's schedule



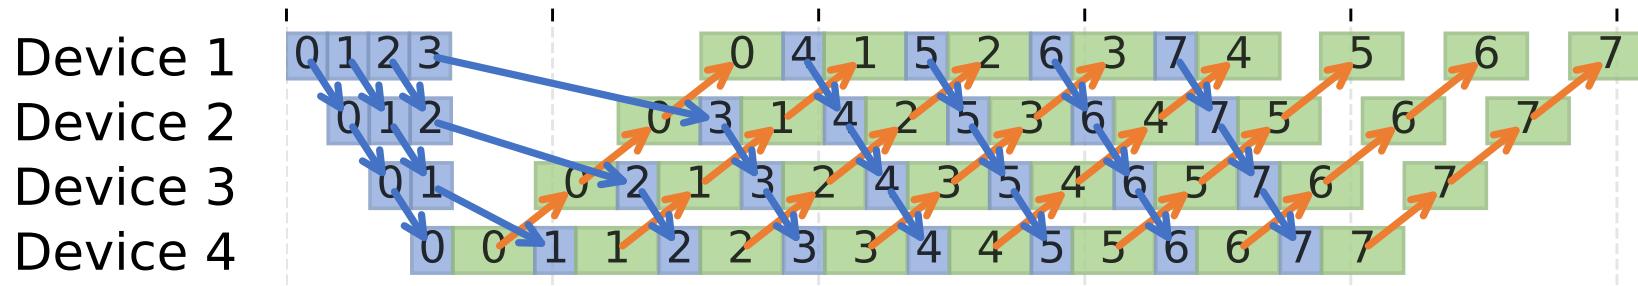
1F1B scheduling is inefficient under dynamic micro-batches.

Our solution: an efficient schedule minimizing GPU idling for pipeline parallel training

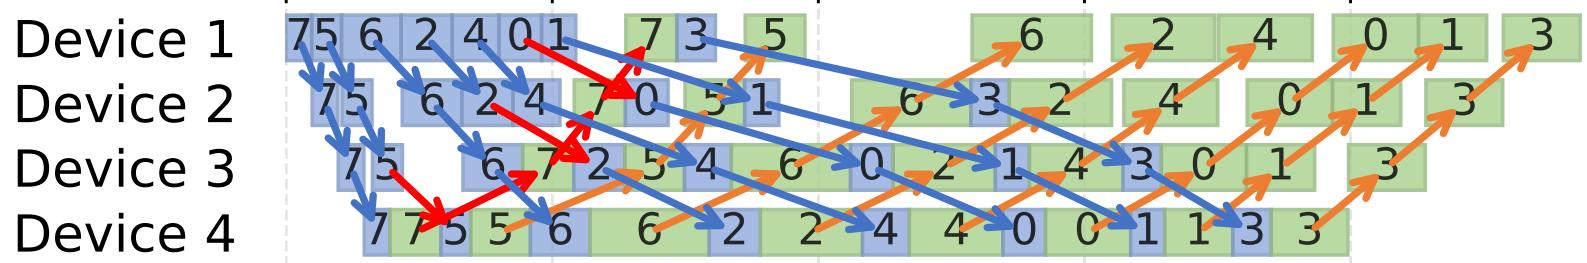
# Motivation

## Challenges of dynamic micro-batching

Improper communication order between pipeline stages may lead to deadlocks in dynamic pipelines.



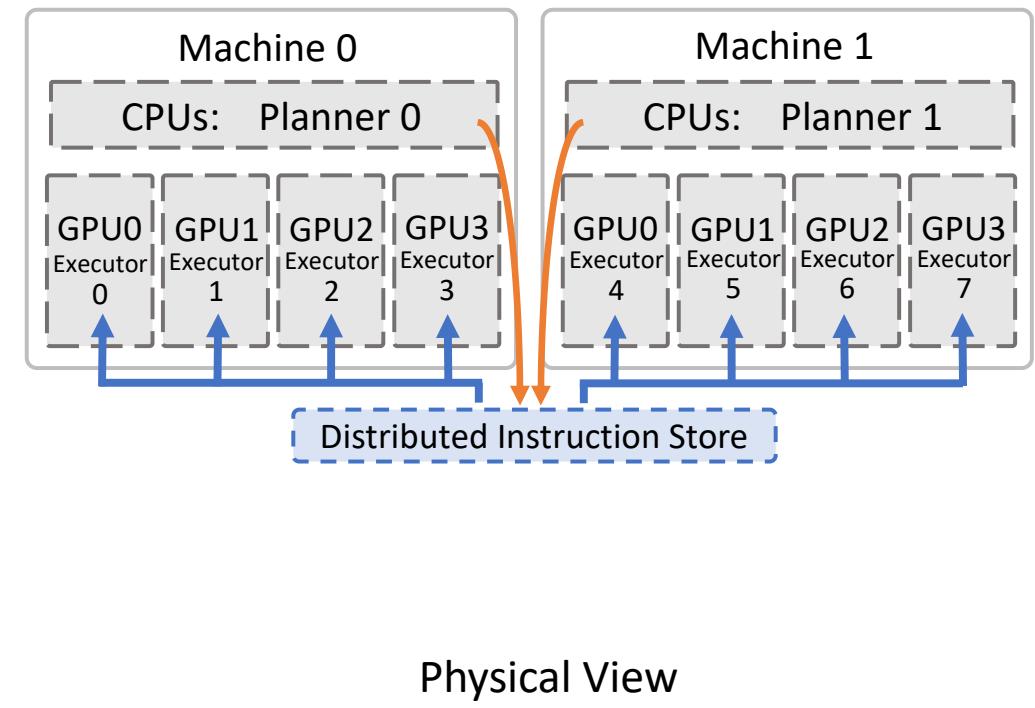
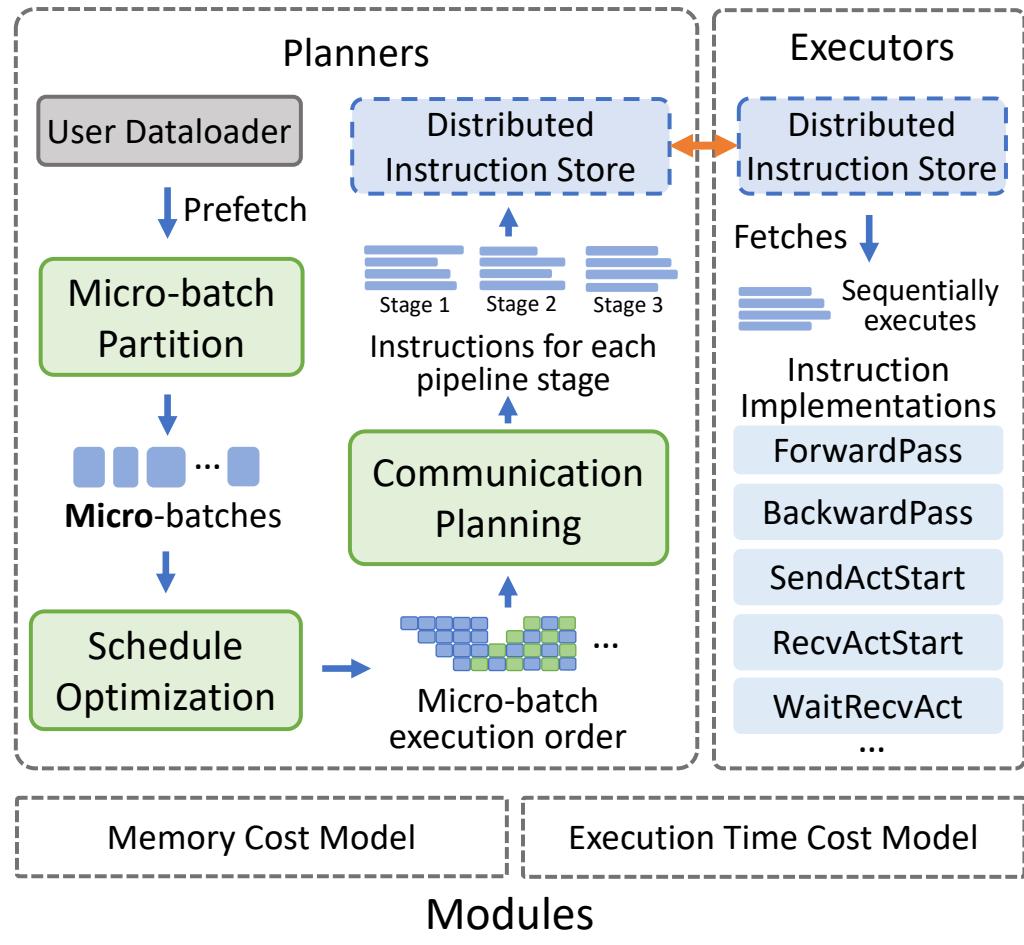
Regular communication pattern in 1F1B schedule



Irregular communication pattern in DynaPipe which changes each iteration

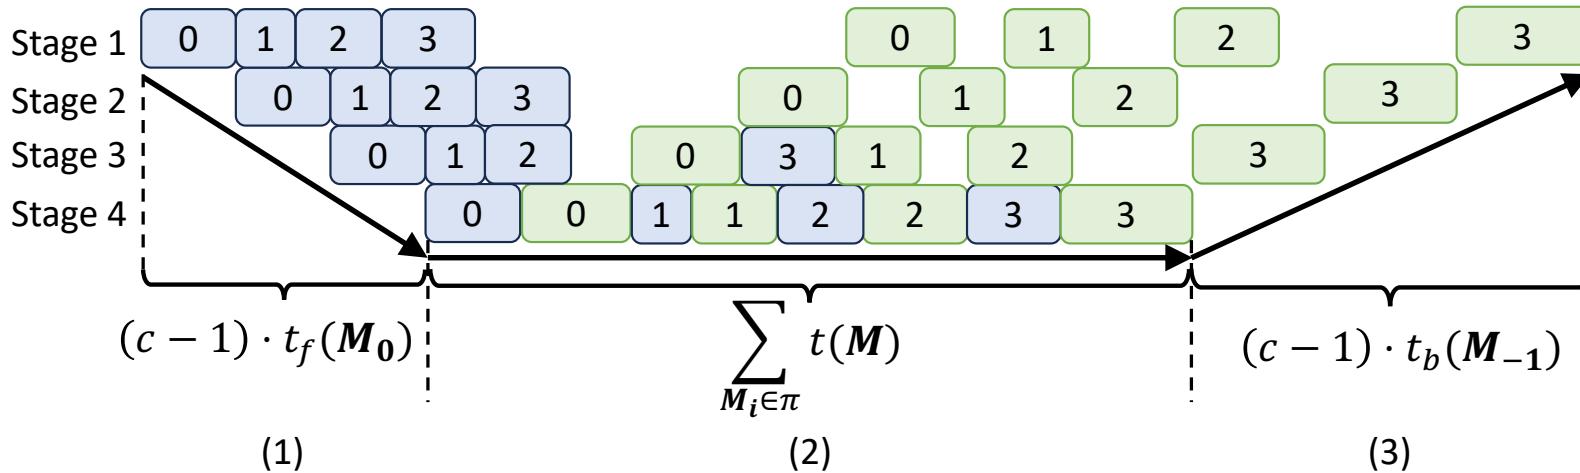
Our solution: plan the communication order for each iteration in advance

# Overview



# Micro-batch Construction

Modelling pipeline execution time



$$\min_{\pi} \left\{ (c - 1) \cdot \max\{t(\mathbf{M}_i) | \mathbf{M}_i \in \pi\} + \sum_{\mathbf{M}_i \in \pi} t(\mathbf{M}_i) \right\}$$

$\pi$ : The set of micro-batches

$c$ : Number of pipeline stages

$t(M_i)$ : execution time of micro-batch  $M_i$

# Micro-batch Construction

DP algorithm

$$f(n; t_{max}) = \min_{1 \leq i \leq n-1} \{ f(i; t_{max}) + t(\mathbf{M}_{S[i+1:n]}) \mid t(\mathbf{M}_{S[i+1:n]}) \leq t_{max} \}$$

$f(n; t_{max})$ : Execution time when optimally partitioning samples 1...n into micro-batches, while each micro-batch's execution time is less than  $t_{max}$ .

$t(M_{s[i:j]})$ : execution time of micro-batch consisting of samples i...j

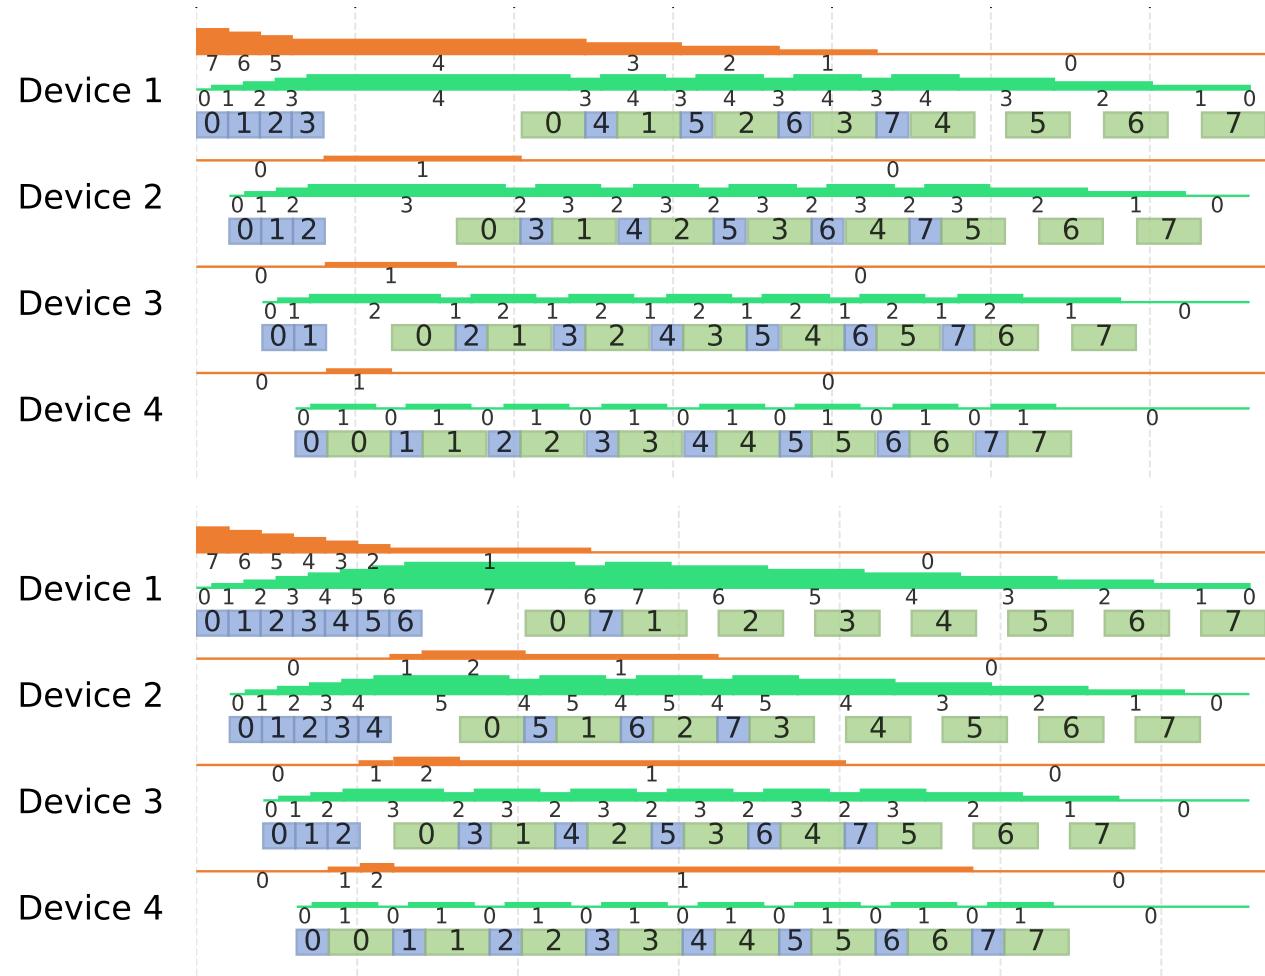
Goal: Find  $\min_{t_{max}} f(N; t_{max})$

# Pipeline Execution Schedule

Adaptive schedule: controlling the injection time of micro-batches

**Safety Stocks:** number of ready (all previous stages have been completed) micro-batches at each device

**1F1B:** 0 safety stocks during steady state, prone to execution time variation

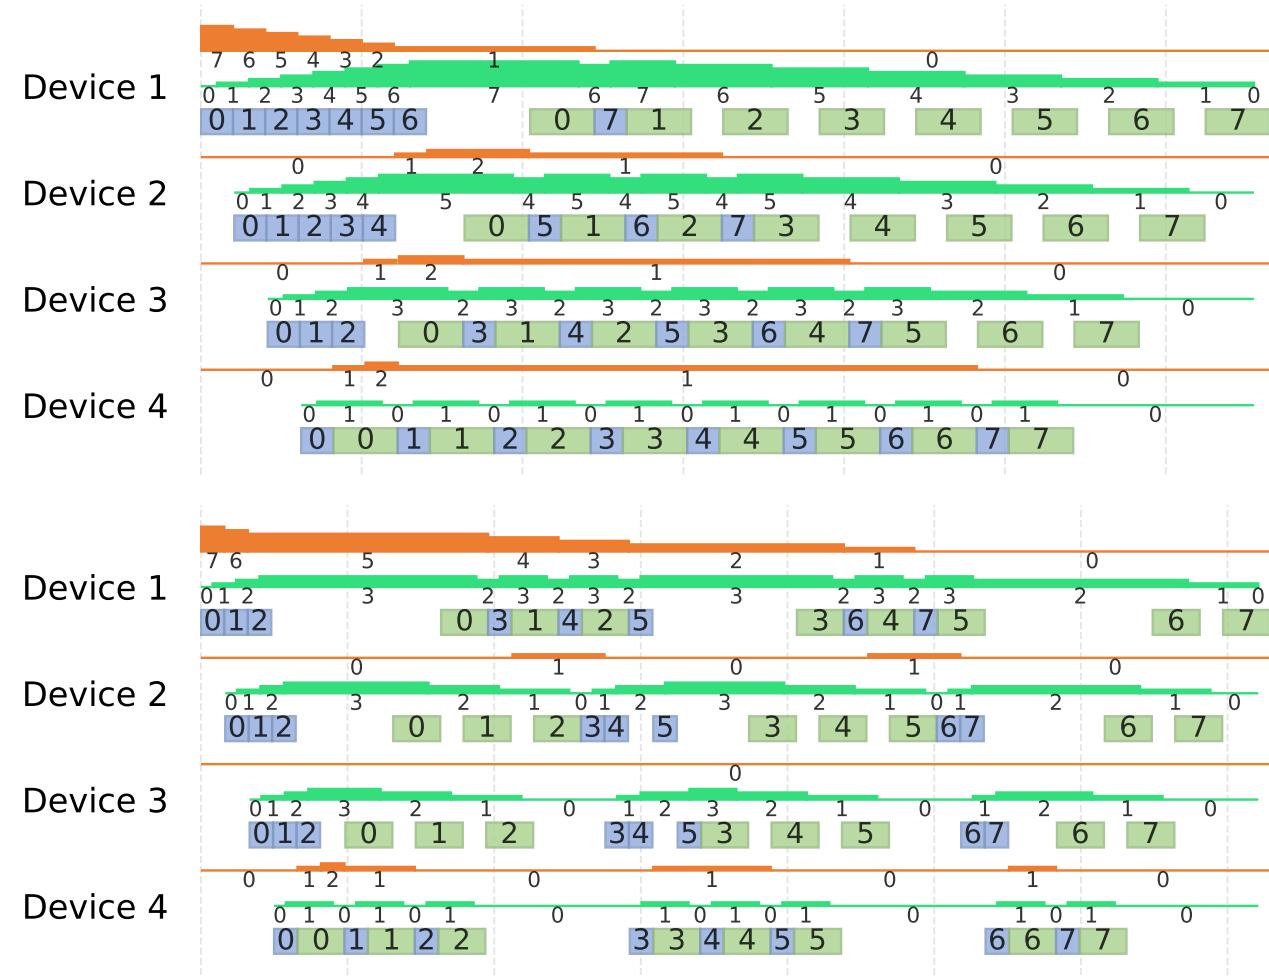


**Inject more micro-batches in the beginning of iteration:** 1 safety stocks during steady state, more robust to variation

# Pipeline Execution Schedule

Adaptive schedule: controlling the injection time of micro-batches

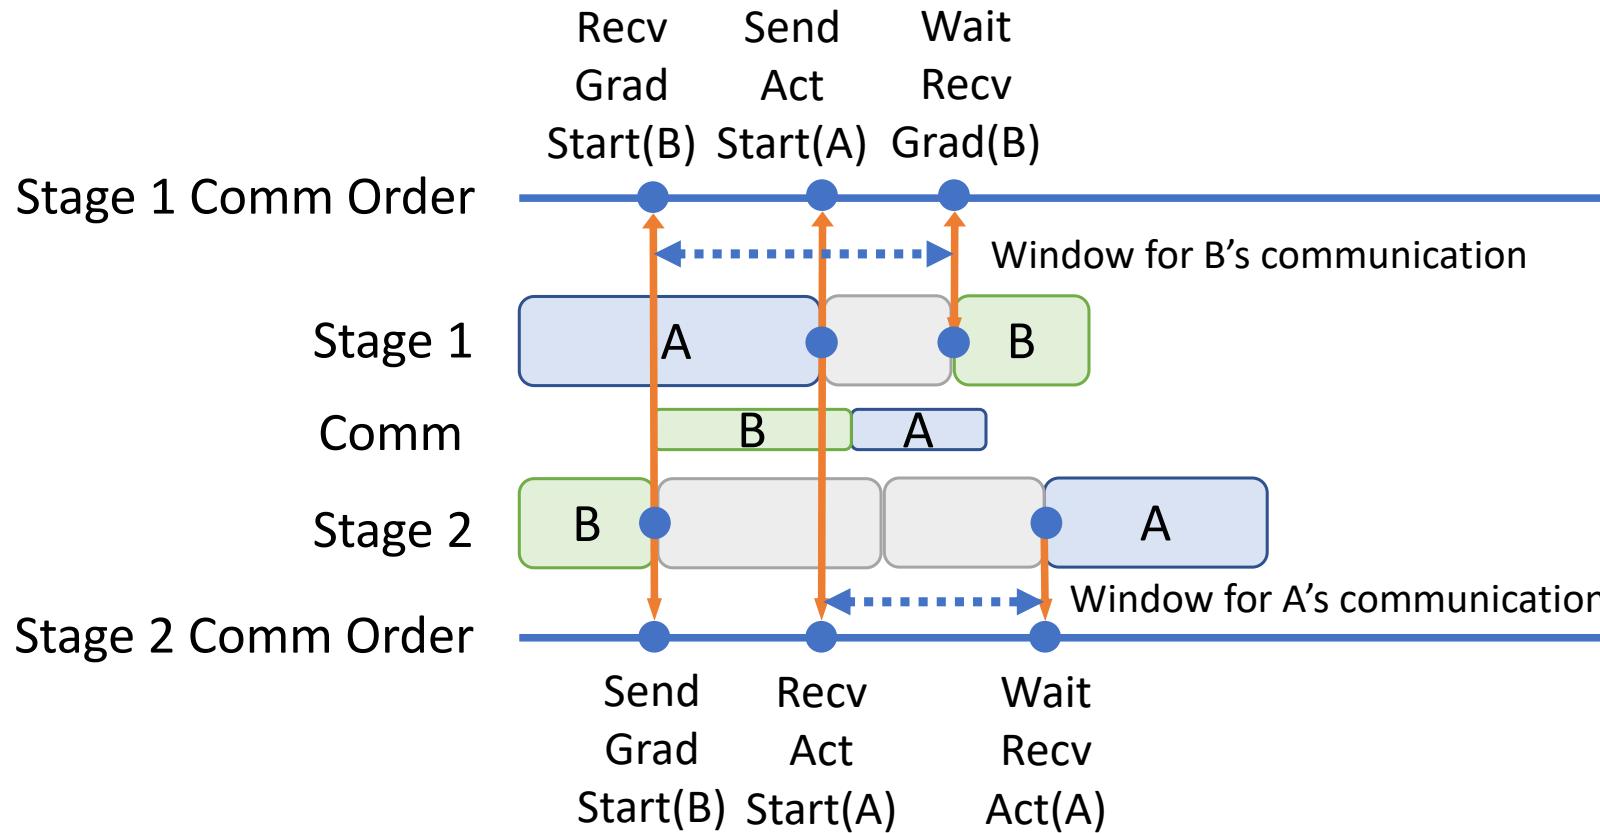
**Inject more micro-batches in the beginning of iteration:** also consumes more memory



**When not enough memory:** delay injection of micro-batches, reducing peak memory from 7 micro-batches' activation to 3

# Communication Planning

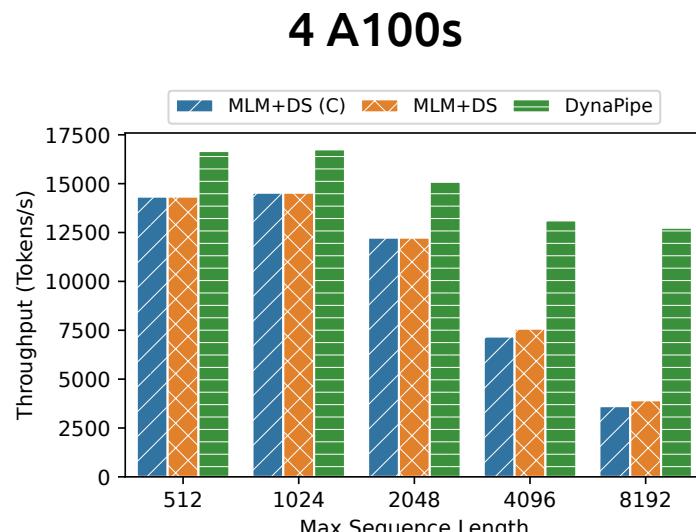
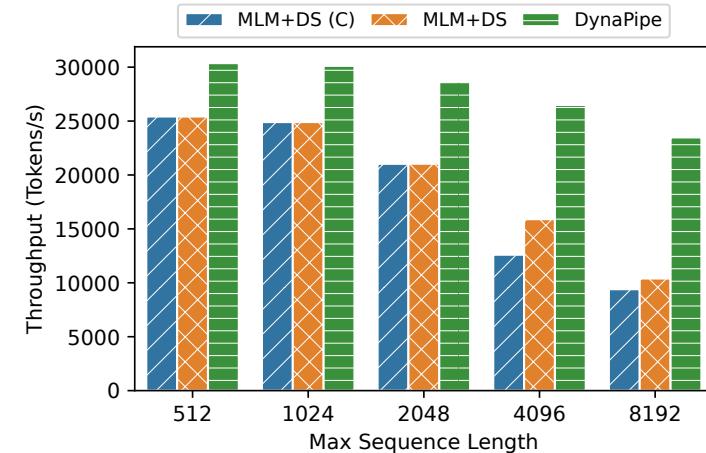
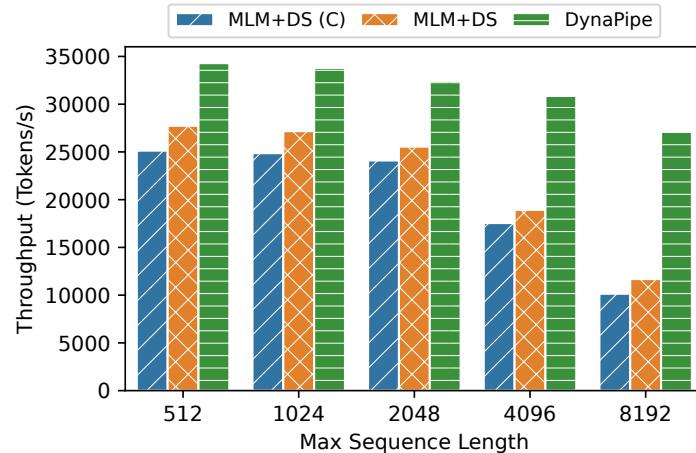
Plan ahead: generate send-recv instruction pairs using simulated timeline



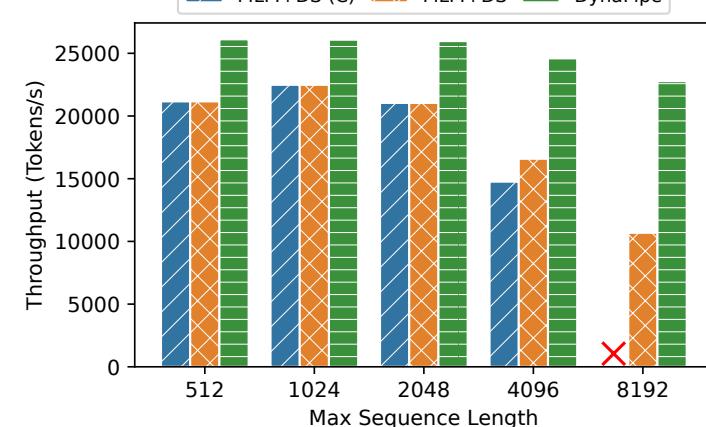
# Evaluation

## GPT on FLANv2 Dataset

Up to **3.25x** speed up  
compared to  
Megatron-LM + DeepSpeed.



**16 A100s**

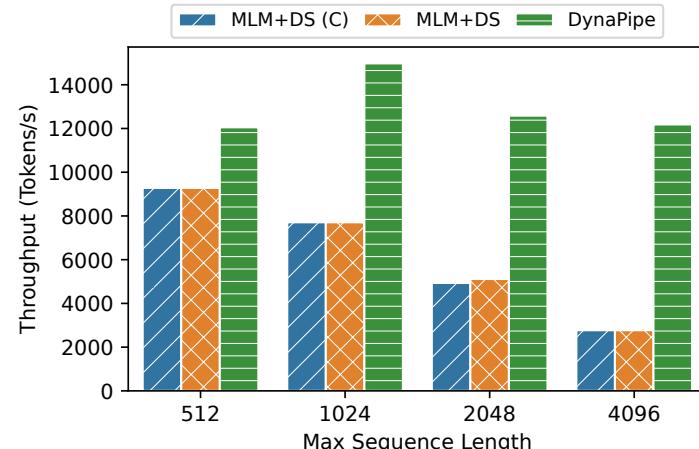


**32 A100s**

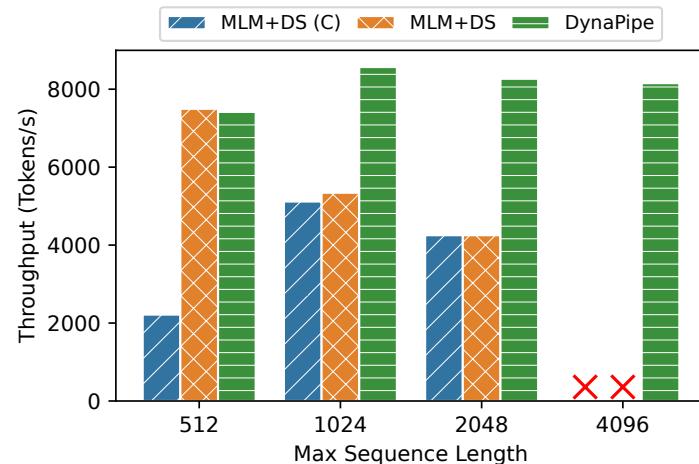
# Evaluation

## T5 on FLANv2 Dataset

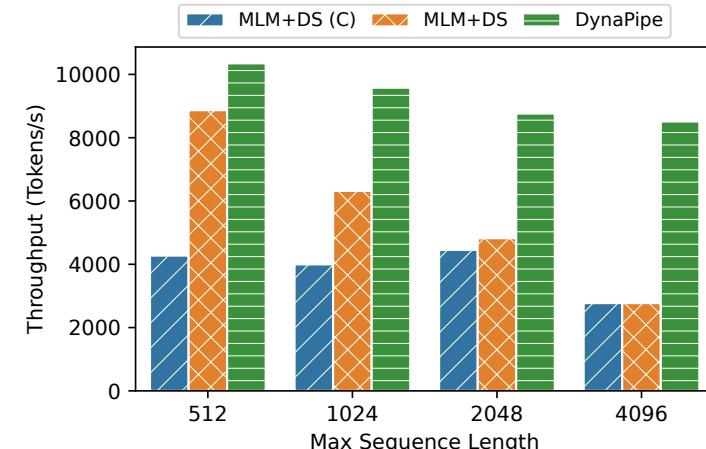
Up to **4.39x** speed up  
compared to  
Megatron-LM + DeepSpeed.



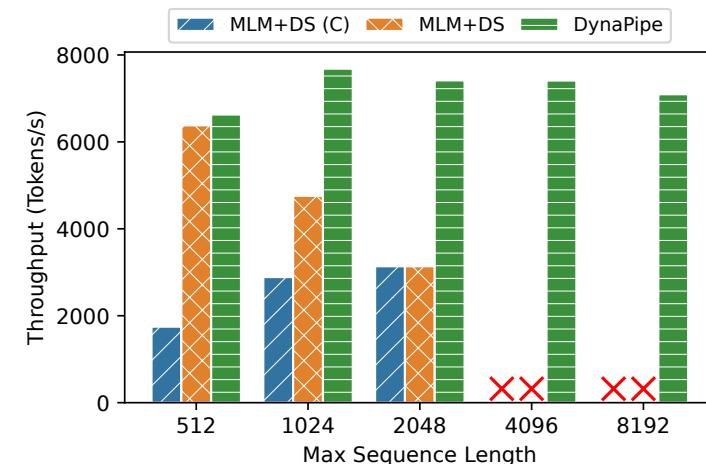
4 A100s



16 A100s



8 A100s

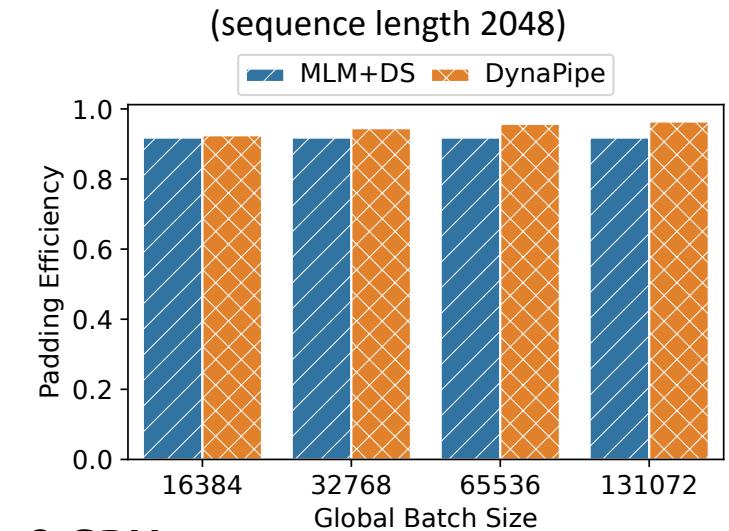
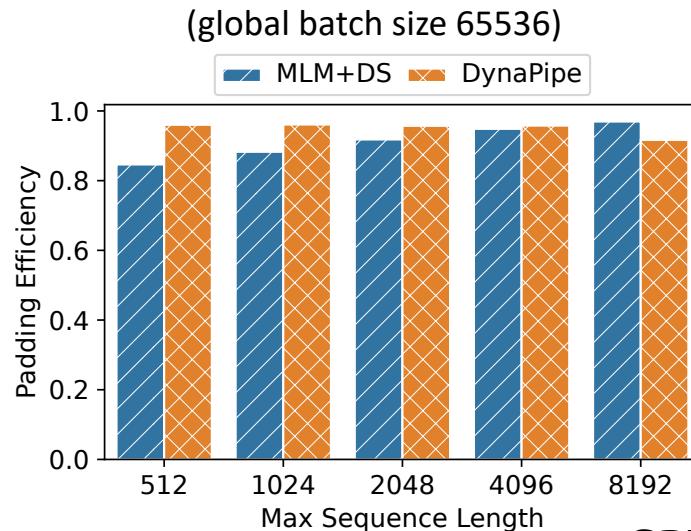


32 A100s

# Evaluation

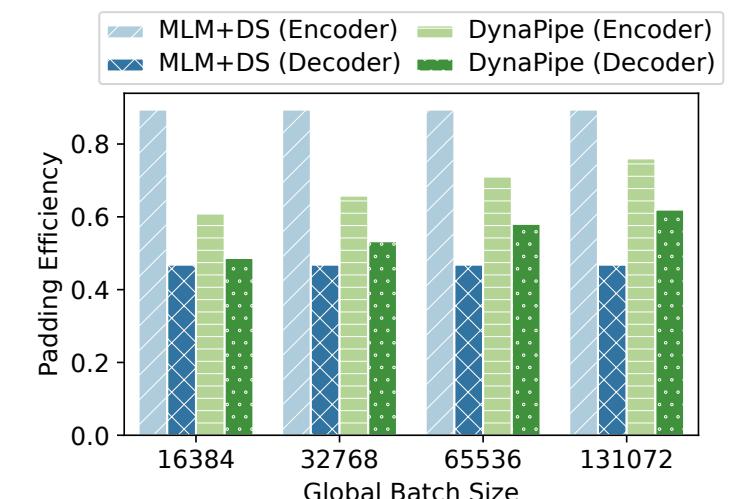
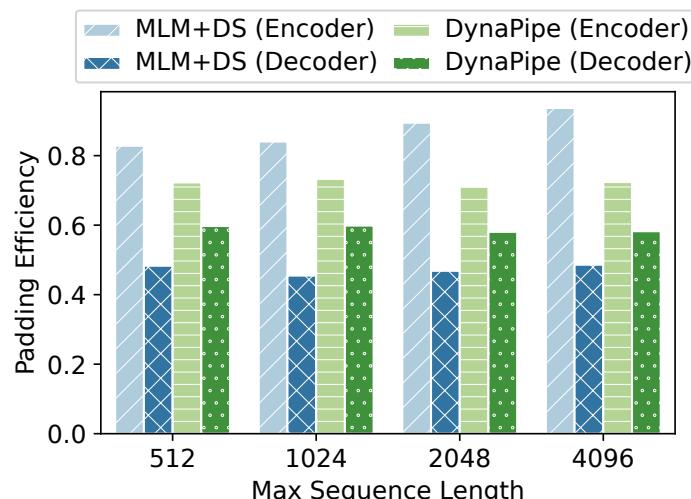
## Padding Efficiency

Comparable padding efficiency as packing for GPT.



GPT on 8 GPUs

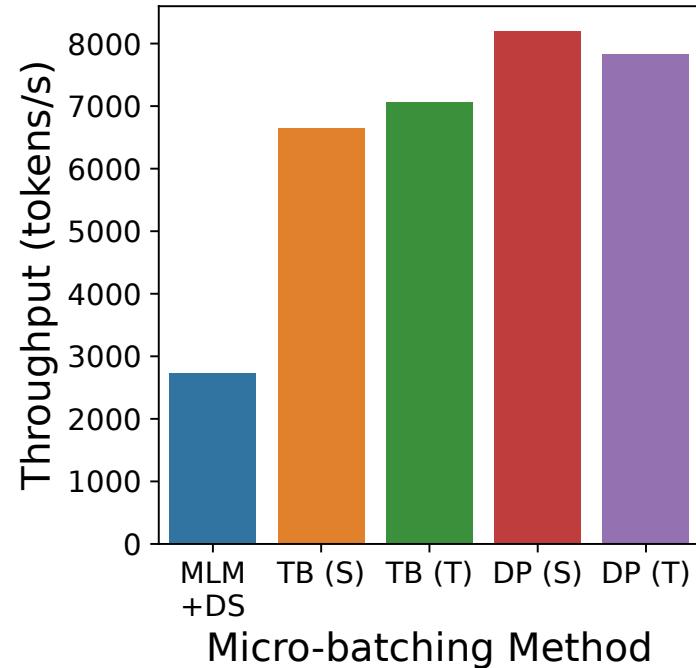
Higher efficiency for T5 decoder, lower for encoder (more balanced overall)



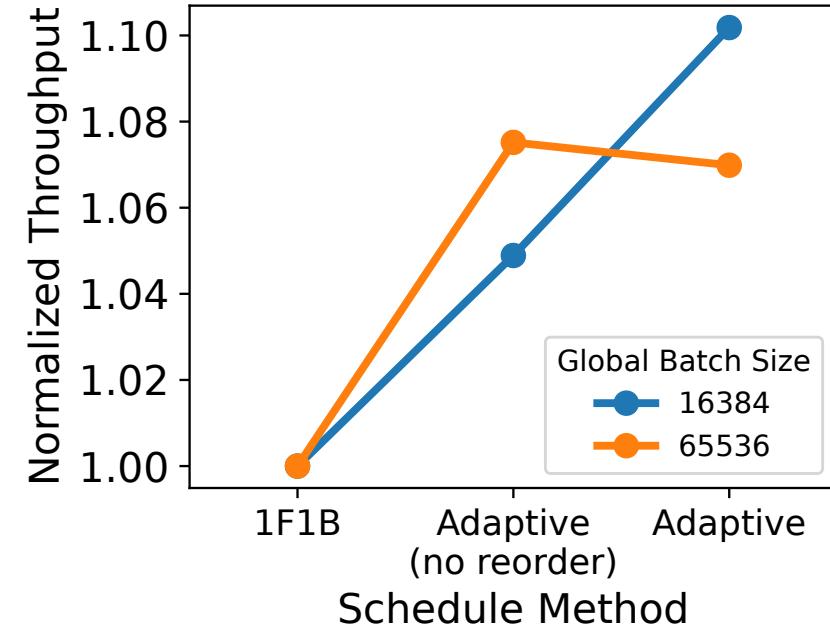
T5 on 8 GPUs

# Evaluation

## Ablation Study



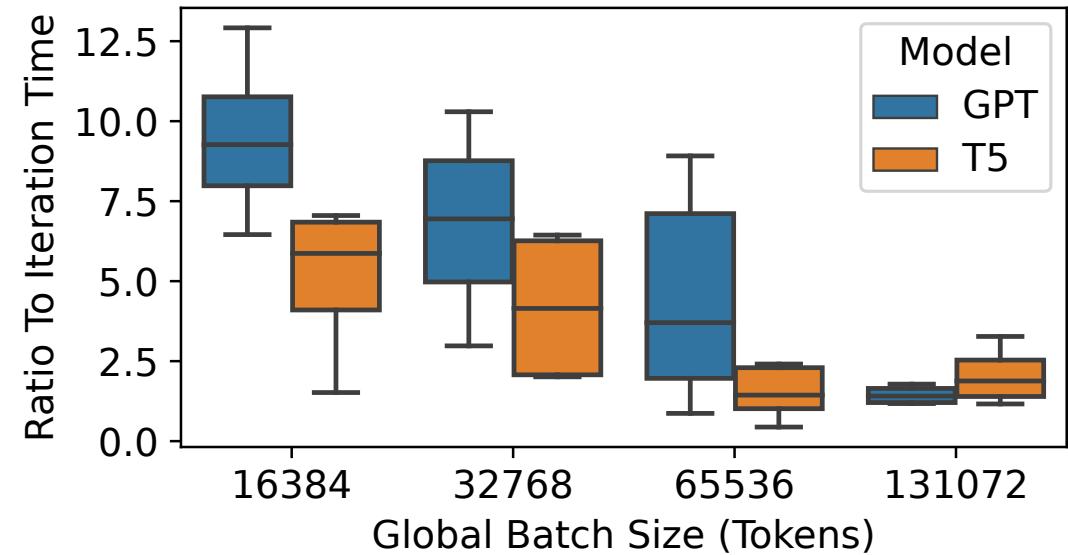
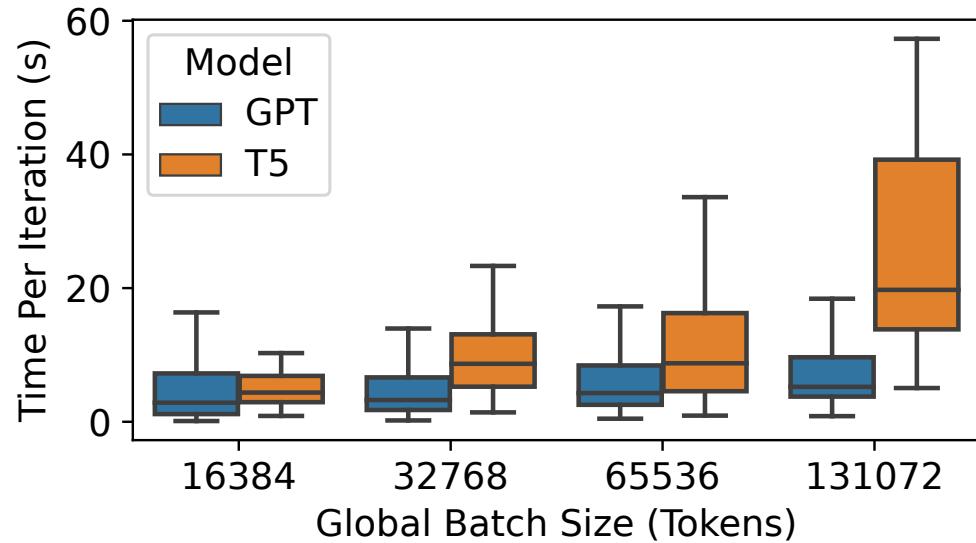
**DP algorithm out-performs best token-based micro-batching methods.**



**Adaptive scheduling out-performs 1F1B.**

# Evaluation

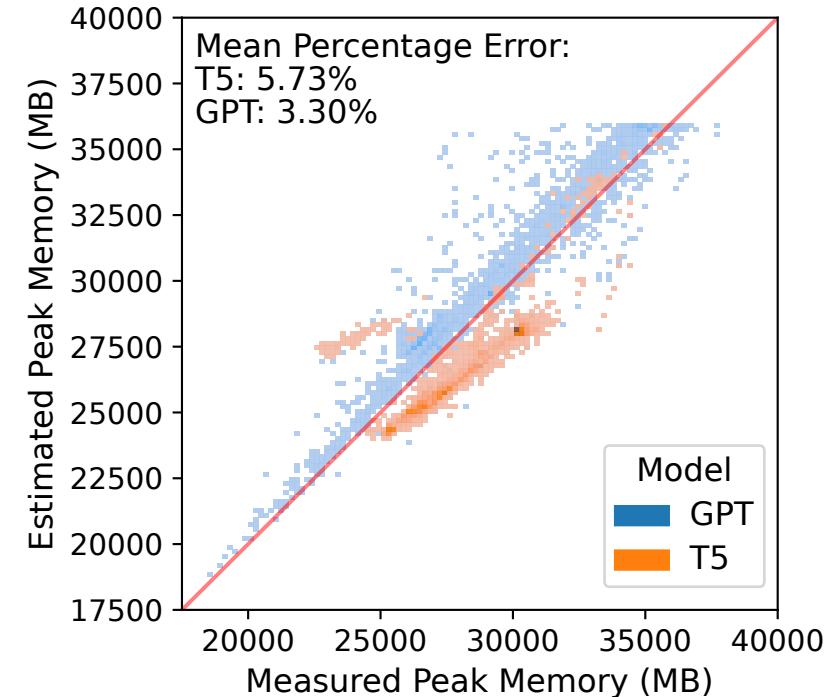
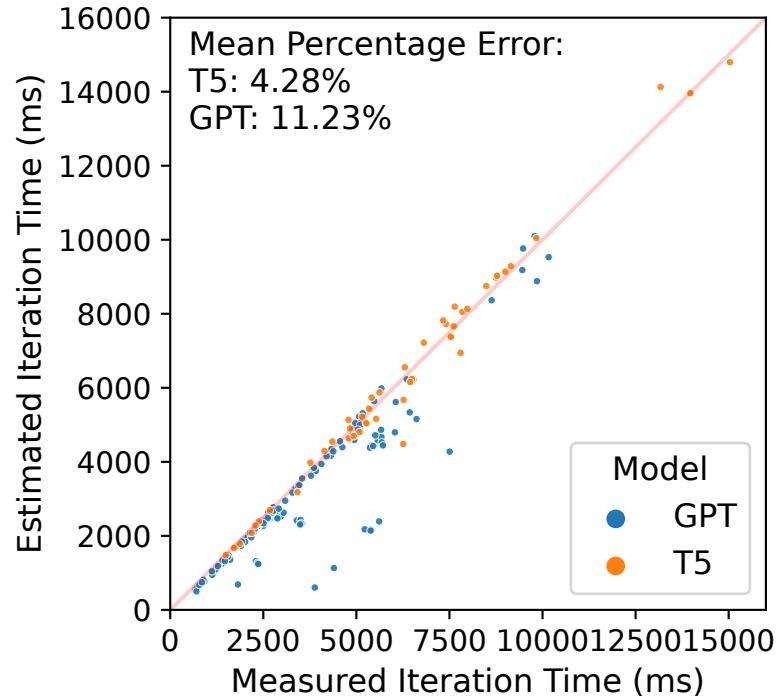
## Planning Time



Planning and model training can fully overlap when parallelized to more than 13 CPU cores in all our experiments.

# Evaluation

## Cost Model Accuracy



Both execution time and memory consumption modelling are accurate, providing useful signal for optimization.